




深度學習

Deep Learning

可解釋性人工智慧
Explainable AI

Instructor: 林英嘉 (Ying-Jia Lin)
2025/05/19



[Course GitHub](#)



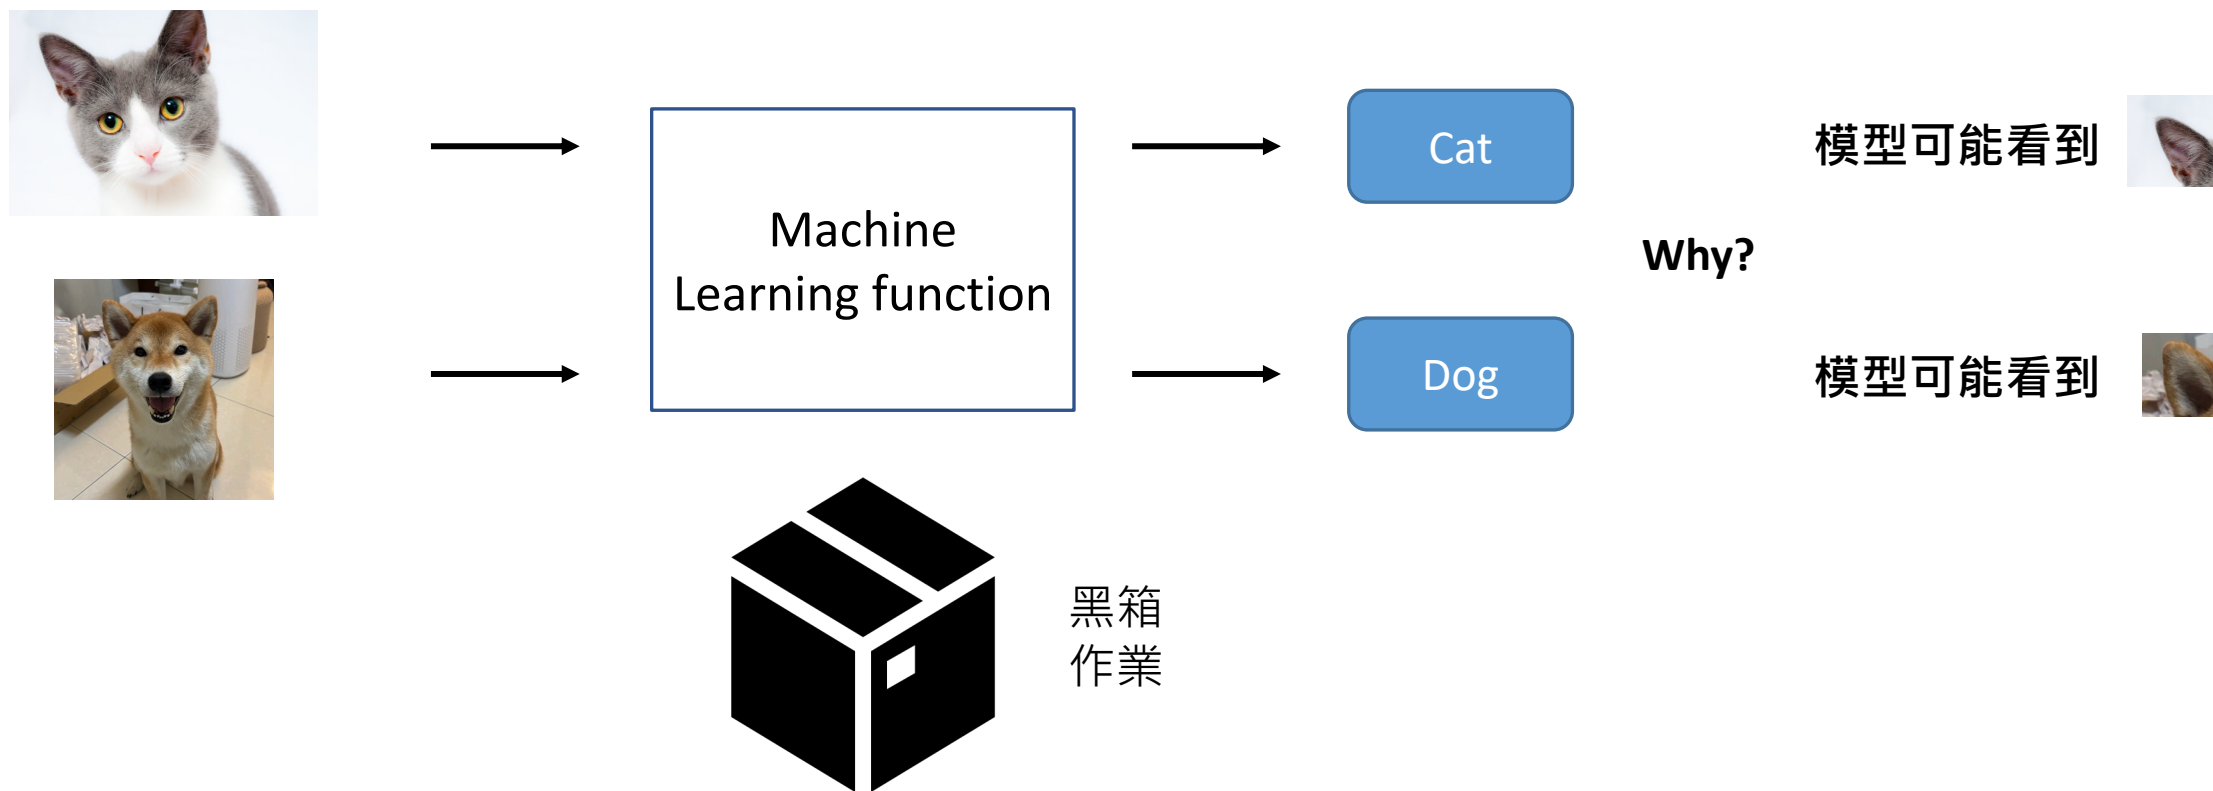
[Slido # DL_0519](#)

Outline

- Introduction
- Class-activation Map (CAM)
- Grad-CAM
- Code



Prediction of a Machine Learning Model




Example: Classification

- Test model: resnet18 pre-trained on ImageNet-1K
- Method: Class activation map (CAM)


Input Image



Pembroke Welsh Corgi



window shade



Visualizing Feature Maps in a CNN

- 151st channel on the conv5 layer of a deep neural network trained on ImageNet




Figure source: Yosinski, Jason, et al. "Understanding neural networks through deep visualization." 2015 ICML Deep Learning Workshop.

Example with SHAP



Figure source:
<https://github.com/shap/shap>



模型的可解釋性與效能


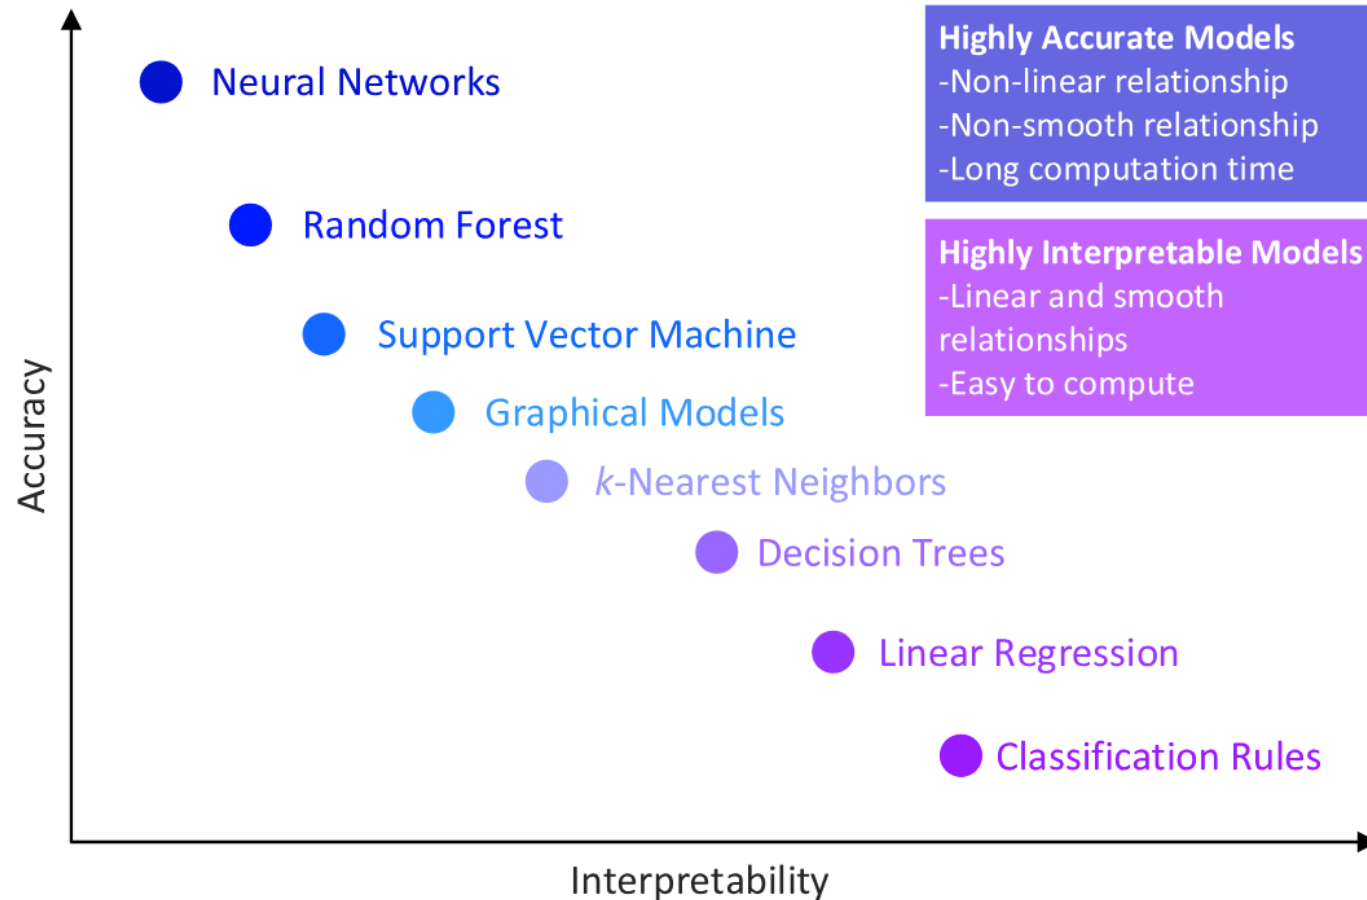


Figure source: Moroch-Cayamcela, Manuel Eugenio, Haeyoung Lee, and Wansu Lim.
"Machine learning for 5G/B5G mobile and wireless communications: Potential, limitations, and future directions." IEEE access 7 (2019): 137184-137206.

Figure source:
<https://9gag.com/gag/aOYA1mE?ref=pn.mw>



Why is Explainable AI important?

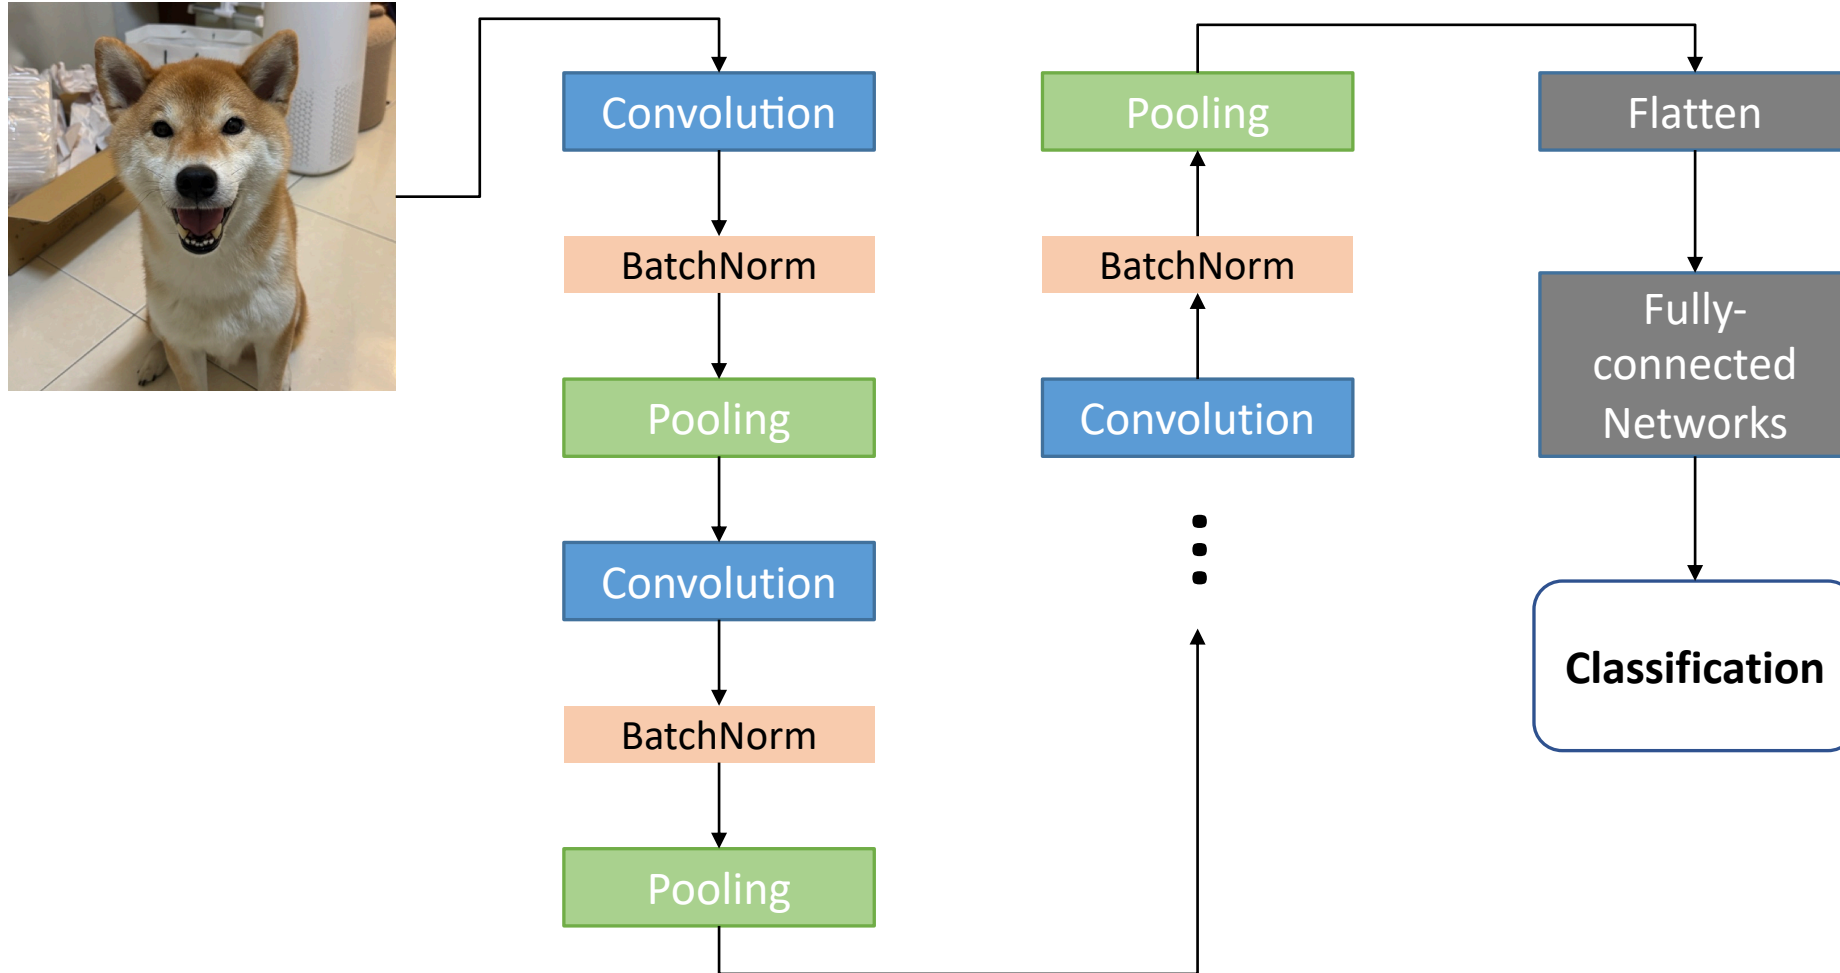
- 確認機器學習模型的判斷合理
 - 建立信任 (使用者 / 政府)
- 改進機器學習模型
 - 從模型輸出找出改進的策略



Convolutional Neural Networks (Recap)

[Recap] The whole Process of a CNN

CNN: Convolutional Neural Networks



[Recap] Convolutions (stride = 1)

1	1	1	1	0	0
0	1	1	0	1	0
0	0	1	1	0	0
0	0	1	1	1	0
0	0	0	1	1	0
0	0	0	0	0	0

Stride = 1

1	0	0
0	1	0
0	0	-1

Filter

Element-wise
multiplication

1



[Recap] Convolutions (stride = 1)

1	1	1	1	0	0
0	1	1	0	1	0
0	0	1	1	0	0
0	0	1	1	1	0
0	0	0	1	1	0
0	0	0	0	0	0

Stride = 1

1	0	0
0	1	0
0	0	-1

Filter

Element-wise
multiplication

1	1	1	2
-1	1	1	0
0	0	1	2
0	0	1	2

feature map



[Recep] 2x2 Pooling (example of Max Pooling)

1	3	1	2
-1	1	1	0
0	1	1	0
0	0	1	2



參數：

- kernel_size=2
- stride = 2

3	2
1	2



[Recep] 2x2 Pooling (example of **Average** Pooling)

1	3	1	2
-1	1	1	0
0	1	1	0
0	0	1	2



參數：

- kernel_size=2
- stride = 2

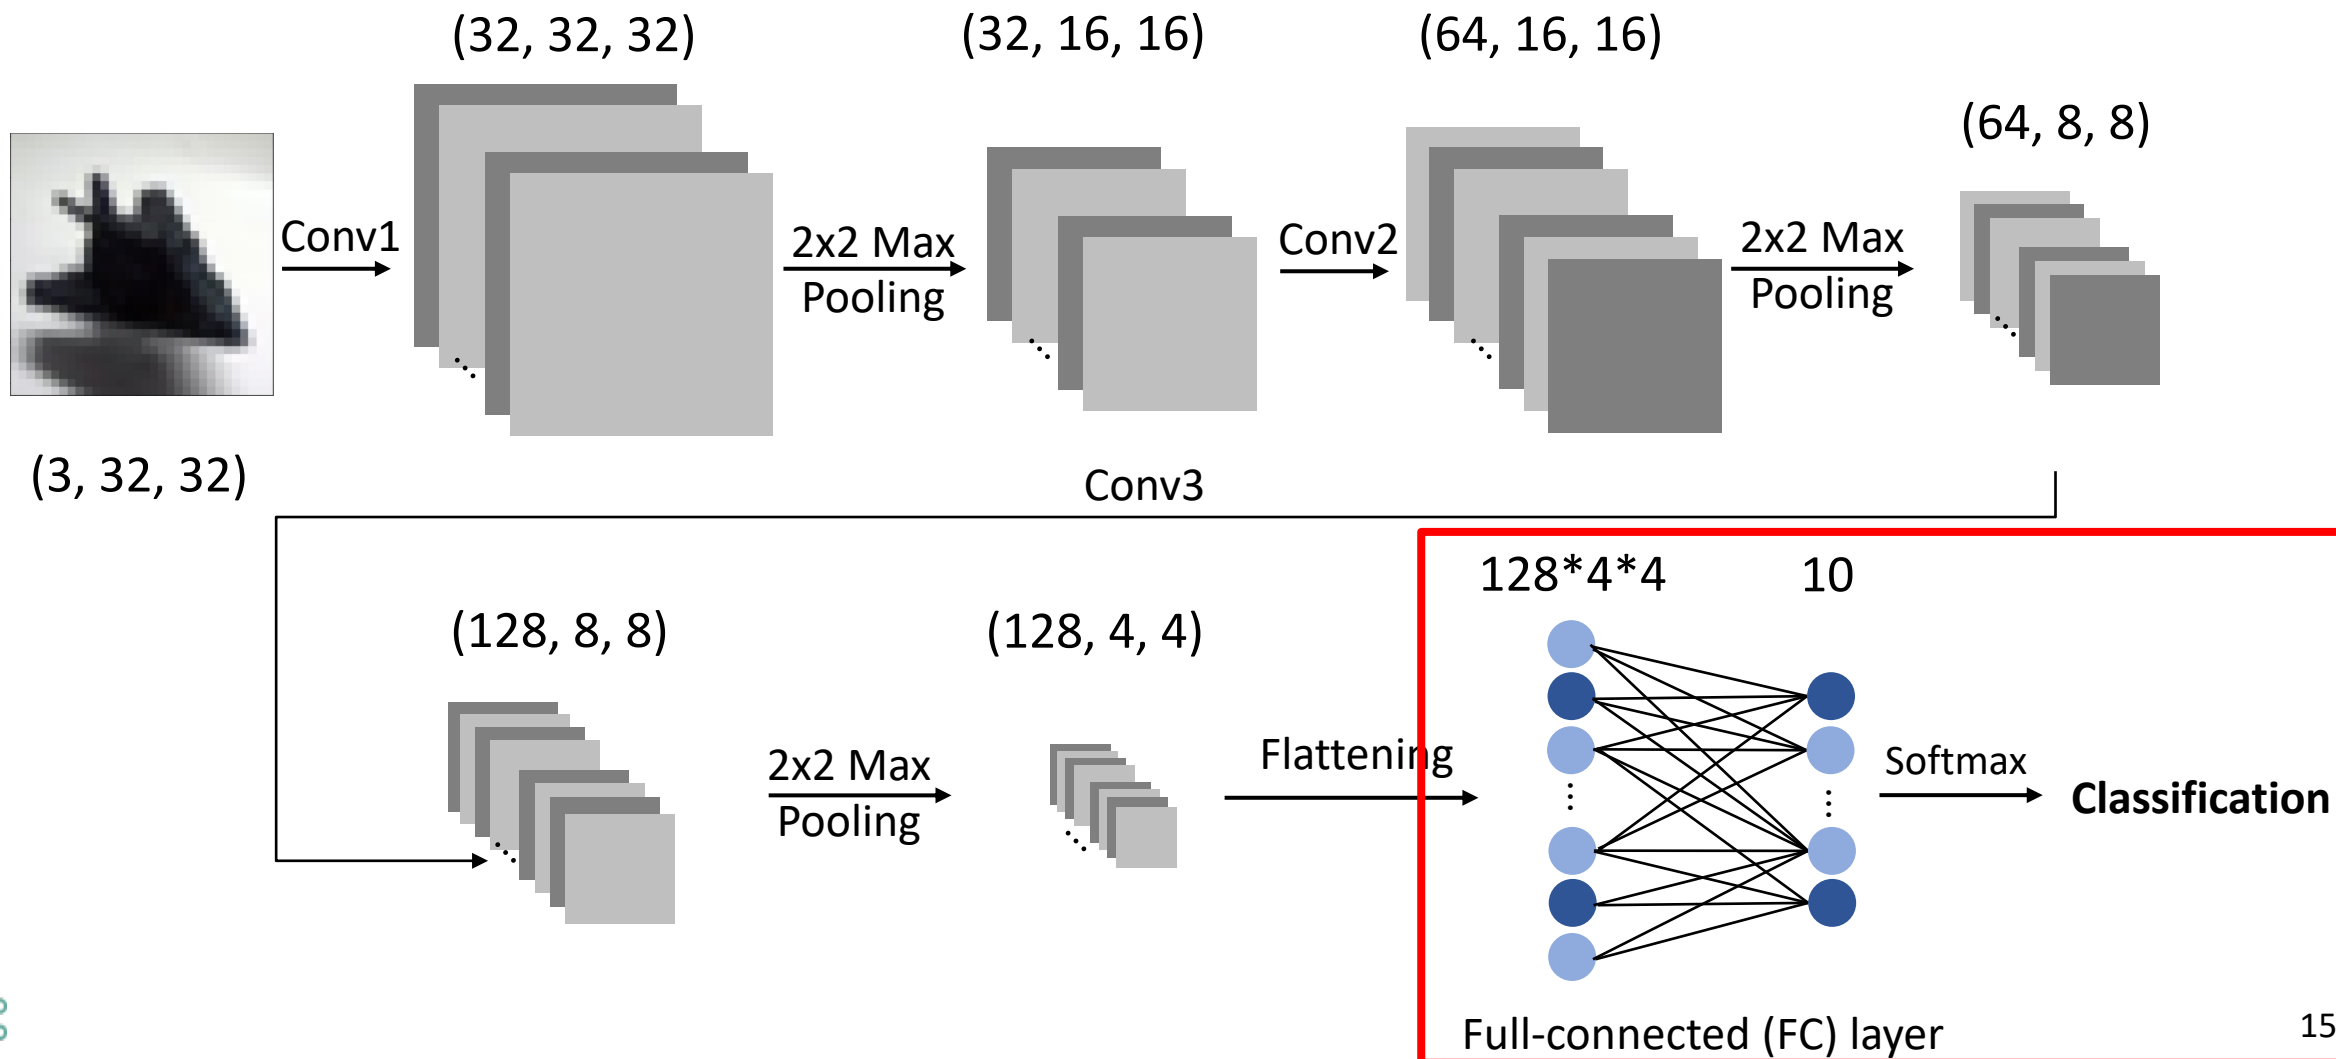
1	1
0.25	1



[Recap] Convolutional Neural Networks (CNN)


*維度意義：(C, H, W)

[CGUDL_2025_Spring/code/pytorch_mnist.ipynb](https://github.com/CGUDL/2025_Spring/code/pytorch_mnist.ipynb)



[Recap] FC layer 參數量

FC: fully-connected



如果不要拉平 (flattening) 呢?

RGB images	參數量比較 (不算 bias 數)
FC layer	$128 * 4 * 4 * 10 = 20480$



Global Average Pooling (GAP)

Feature Map

1	3	1	2
-1	1	1	0
0	1	1	0
0	0	1	2

全局數值取平均

(13/16)

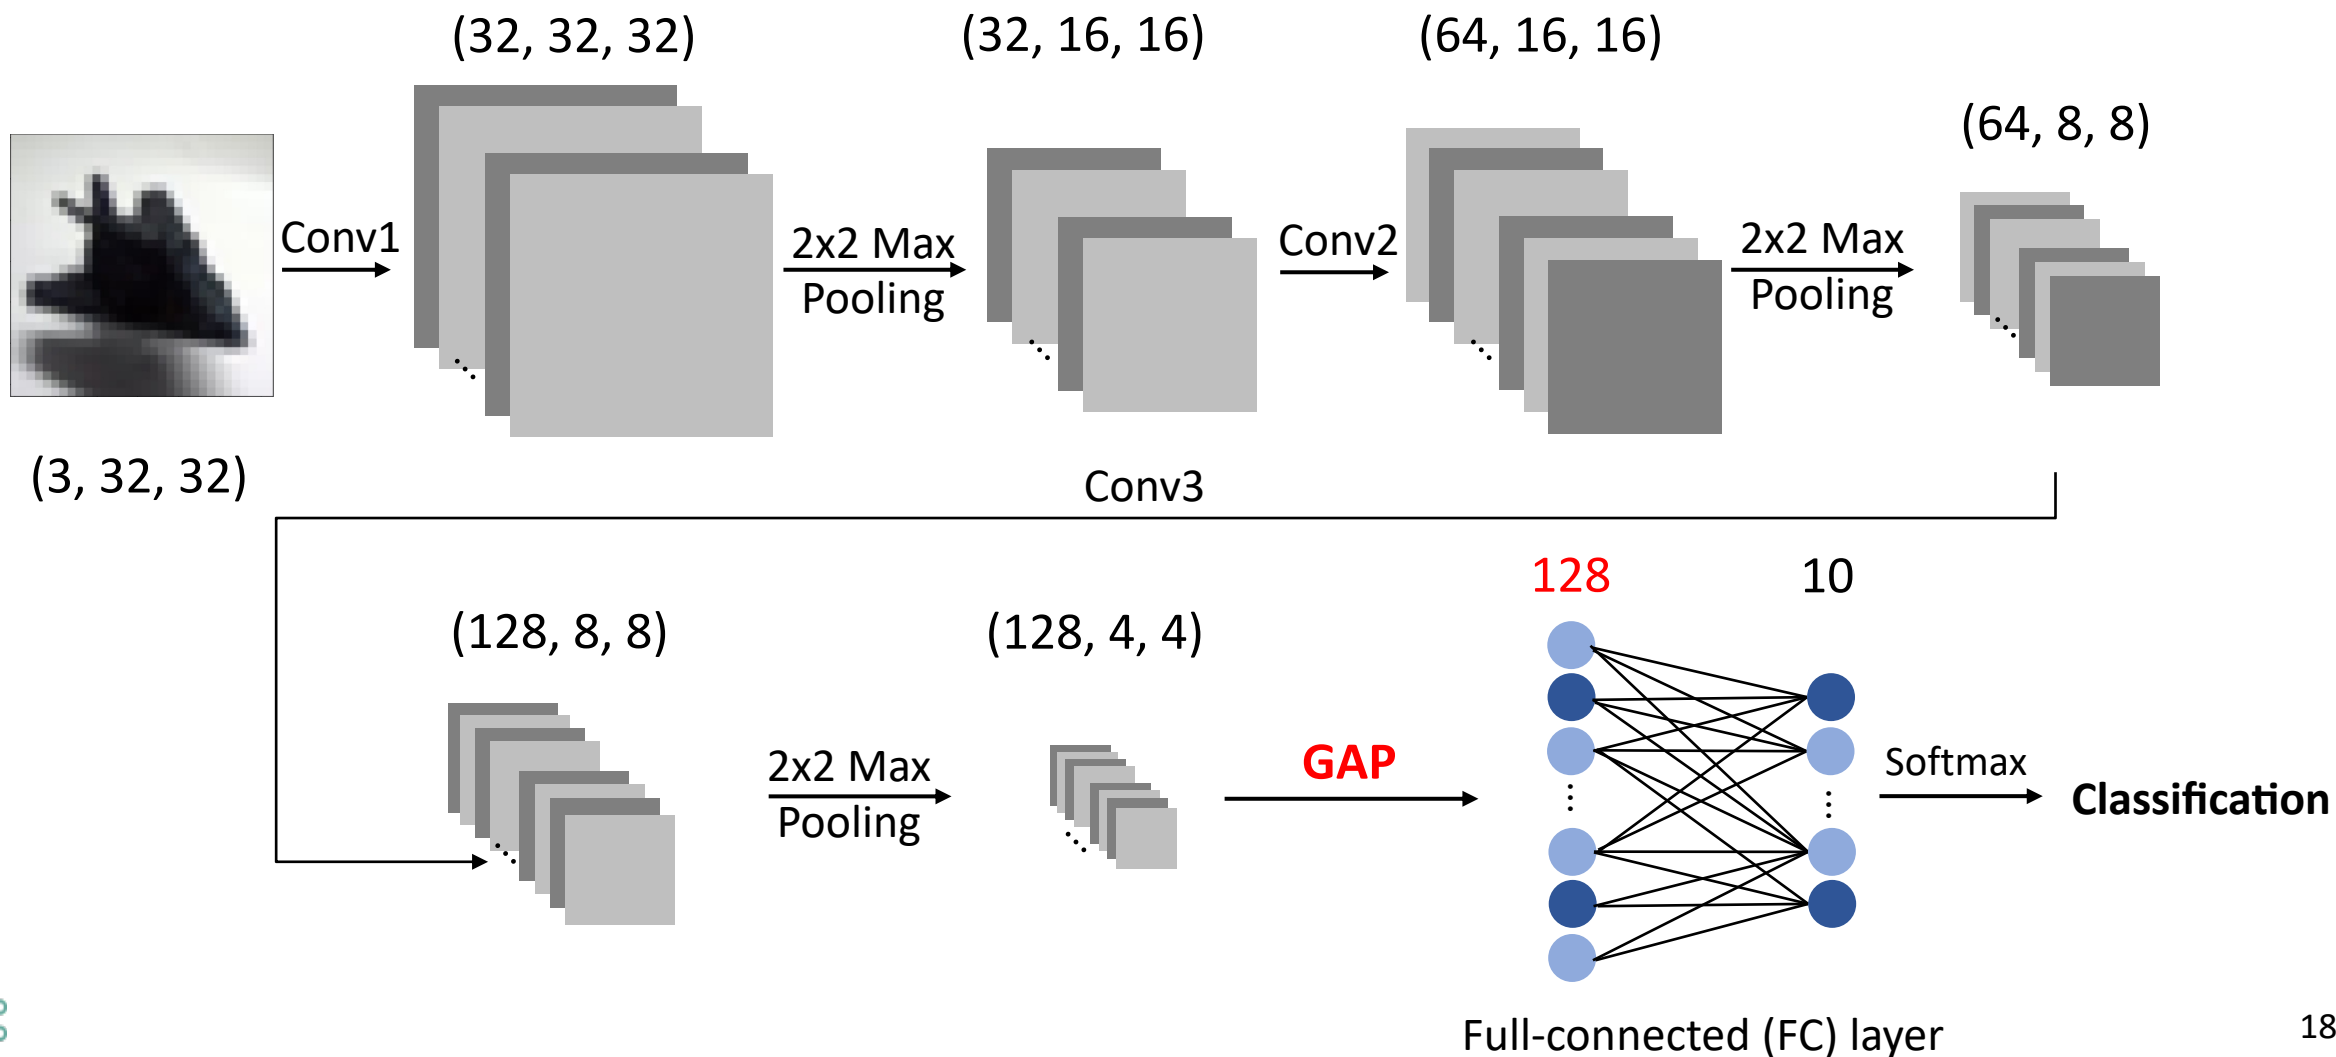
0.8125

一張 feature map
經過GAP後變成一個數值



CNN with Global Average Pooling (GAP)


*維度意義：(C, H, W)




加入 GAP 後參數量下降

FC: fully-connected

Original version



GAP version



RGB images	參數量比較 (不算 bias 數)
Original	$128 * 4 * 4 * 10 = 20480$
GAP	$128 * 10 = 1280$



加入 GAP 後 Testing Error 下降

Table 5: Global average pooling compared to fully connected layer.

	Method	Testing Error
Original	mlpconv + Fully Connected	11.59%
	mlpconv + Fully Connected + Dropout	10.88%
	mlpconv + Global Average Pooling	10.41%



*Dropout 也是減少FC layer中node連接數量的方法



ResNet 架構

<https://arxiv.org/abs/1512.03385>

(Recap)



ResNet 的最後也是 GAP + FC



Class-activation Map (CAM)

[Recap] Convolutional Neural Networks (CNN)

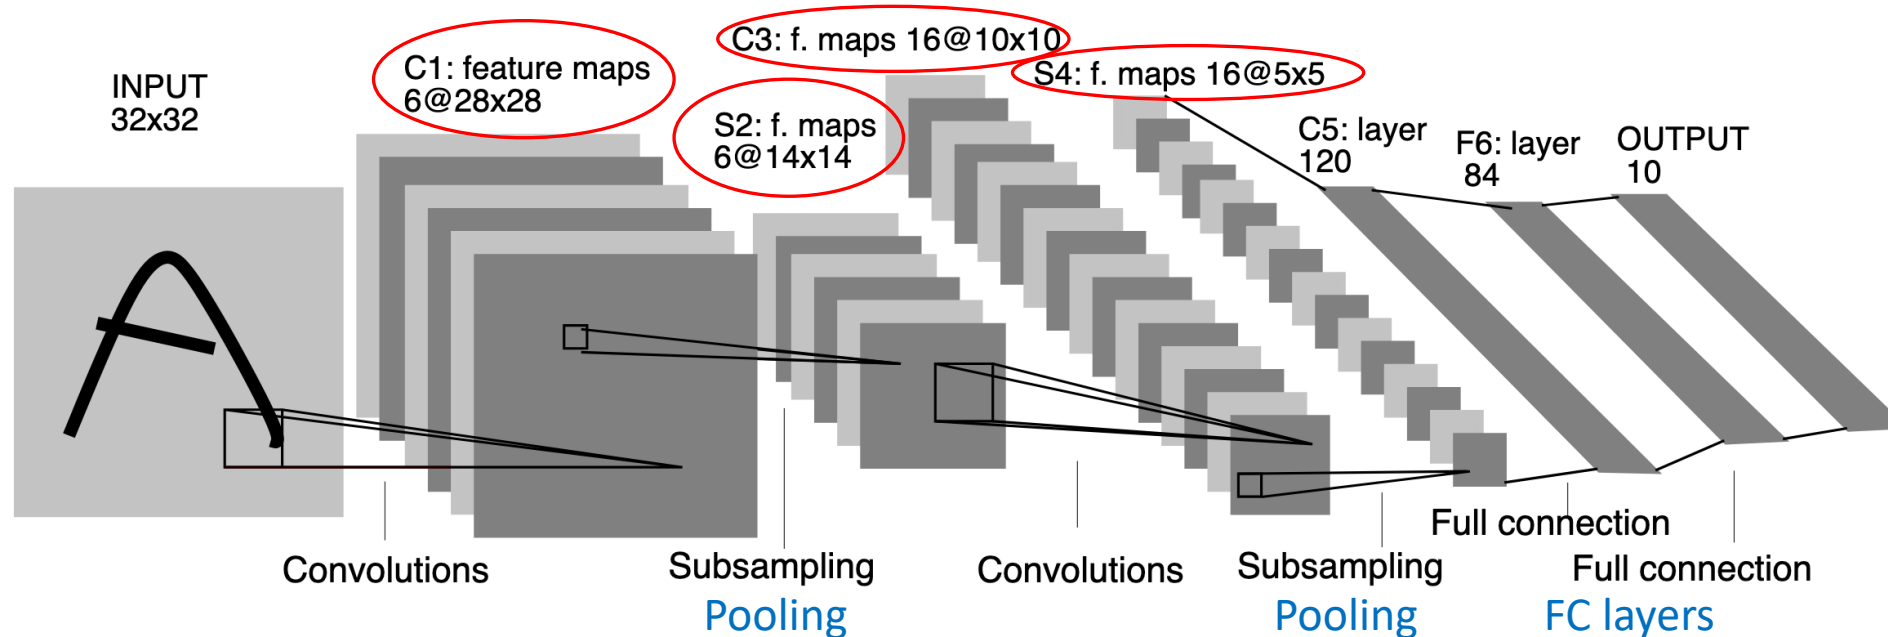


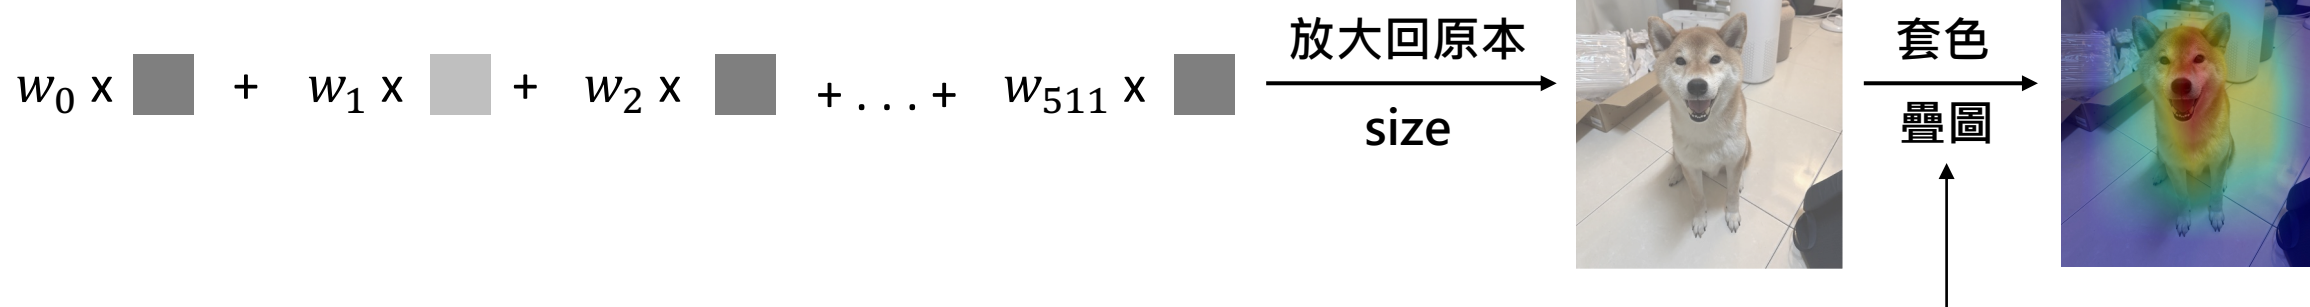
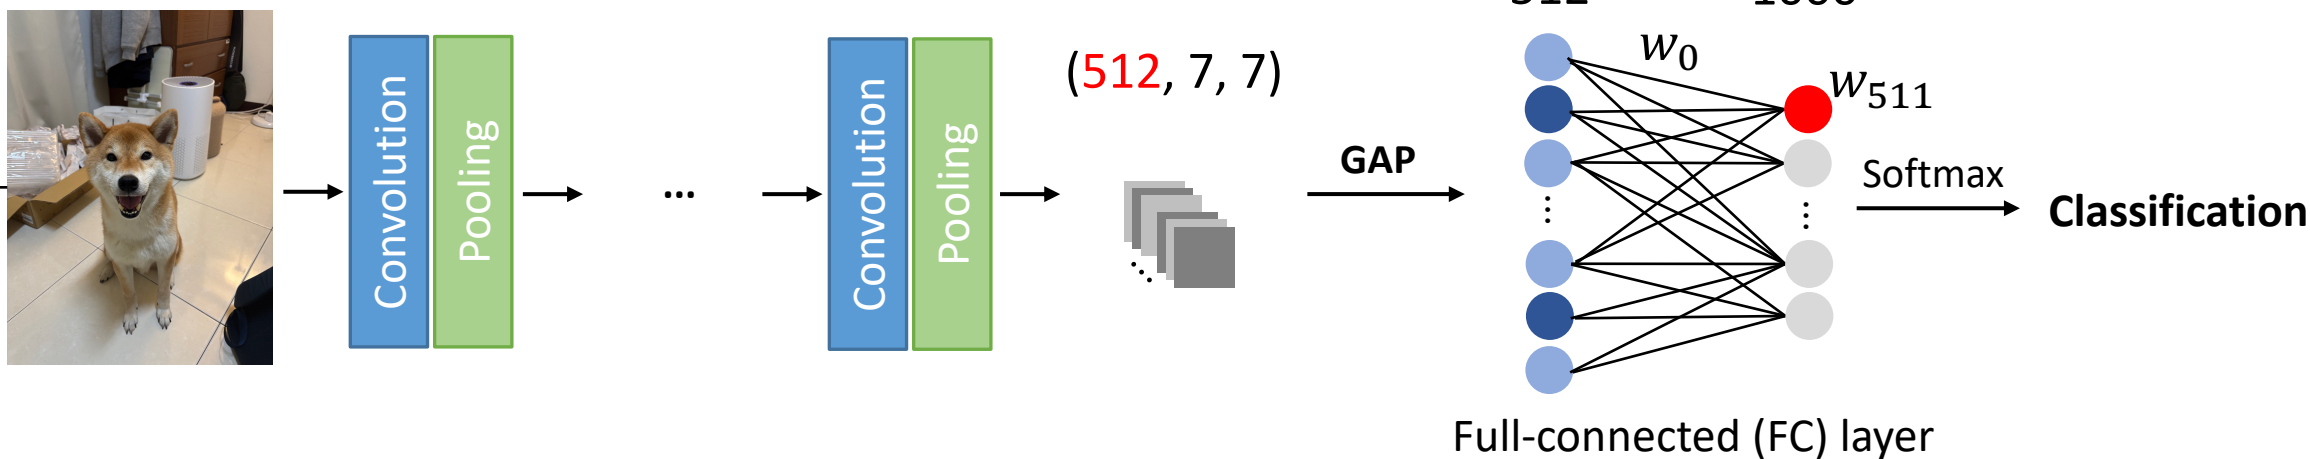
Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278-2324.



Class-activation Map (CAM)

*假設使用 ImageNet pre-trained model



Why Global “Average” Pooling?

How about Global “**Summation**” Pooling?

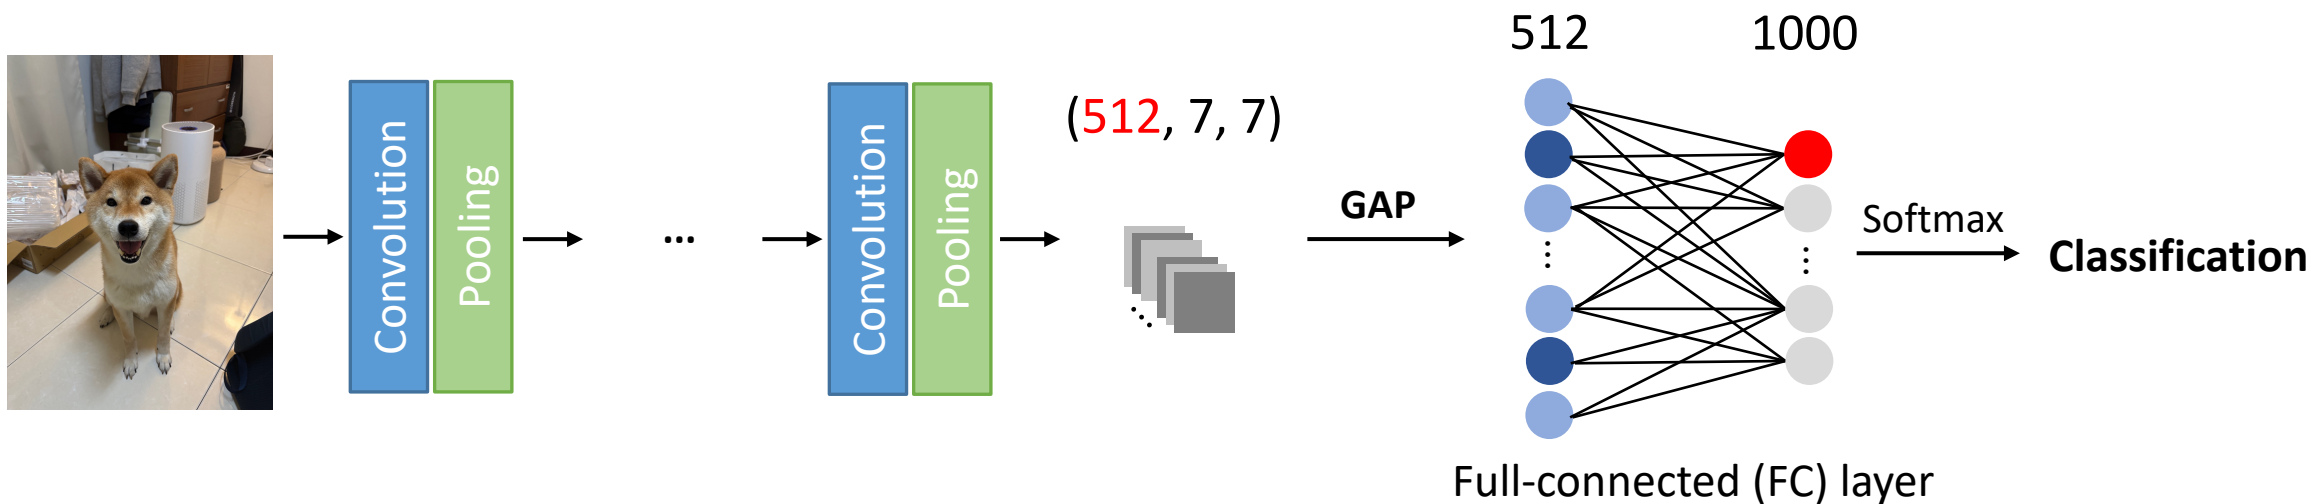
How about Global “**Max**” Pooling (GMP)?

Table 2. Localization error on the ILSVRC validation set. *Backprop* refers to using [22] for localization instead of CAM.

Method	top-1 val.error	top-5 val. error
GoogLeNet-GAP	56.40	43.00
VGGnet-GAP	57.20	45.14
GoogLeNet	60.09	49.34
AlexNet*-GAP	63.75	49.53
AlexNet-GAP	67.19	52.16
NIN	65.47	54.19
Backprop on GoogLeNet	61.31	50.55
Backprop on VGGnet	61.12	51.46
Backprop on AlexNet	65.17	52.64
GoogLeNet-GMP	57.78	45.26



CAM 的問題



不是每個模型最後面都是 GAP + FC
(E.g., VGG-16 的最後是 3 層 FC、ViT 的最後只有 FC)



Grad-CAM (1)

Grad: gradients




能夠產生Feature maps的最後一層



Feature maps
 $A_{i,j}^k$



1000



分類層

i, j : feature map 中x軸與y軸位置

A : feature map

k : feature map 的數目

y^c : 對應到 class c 的 label ID

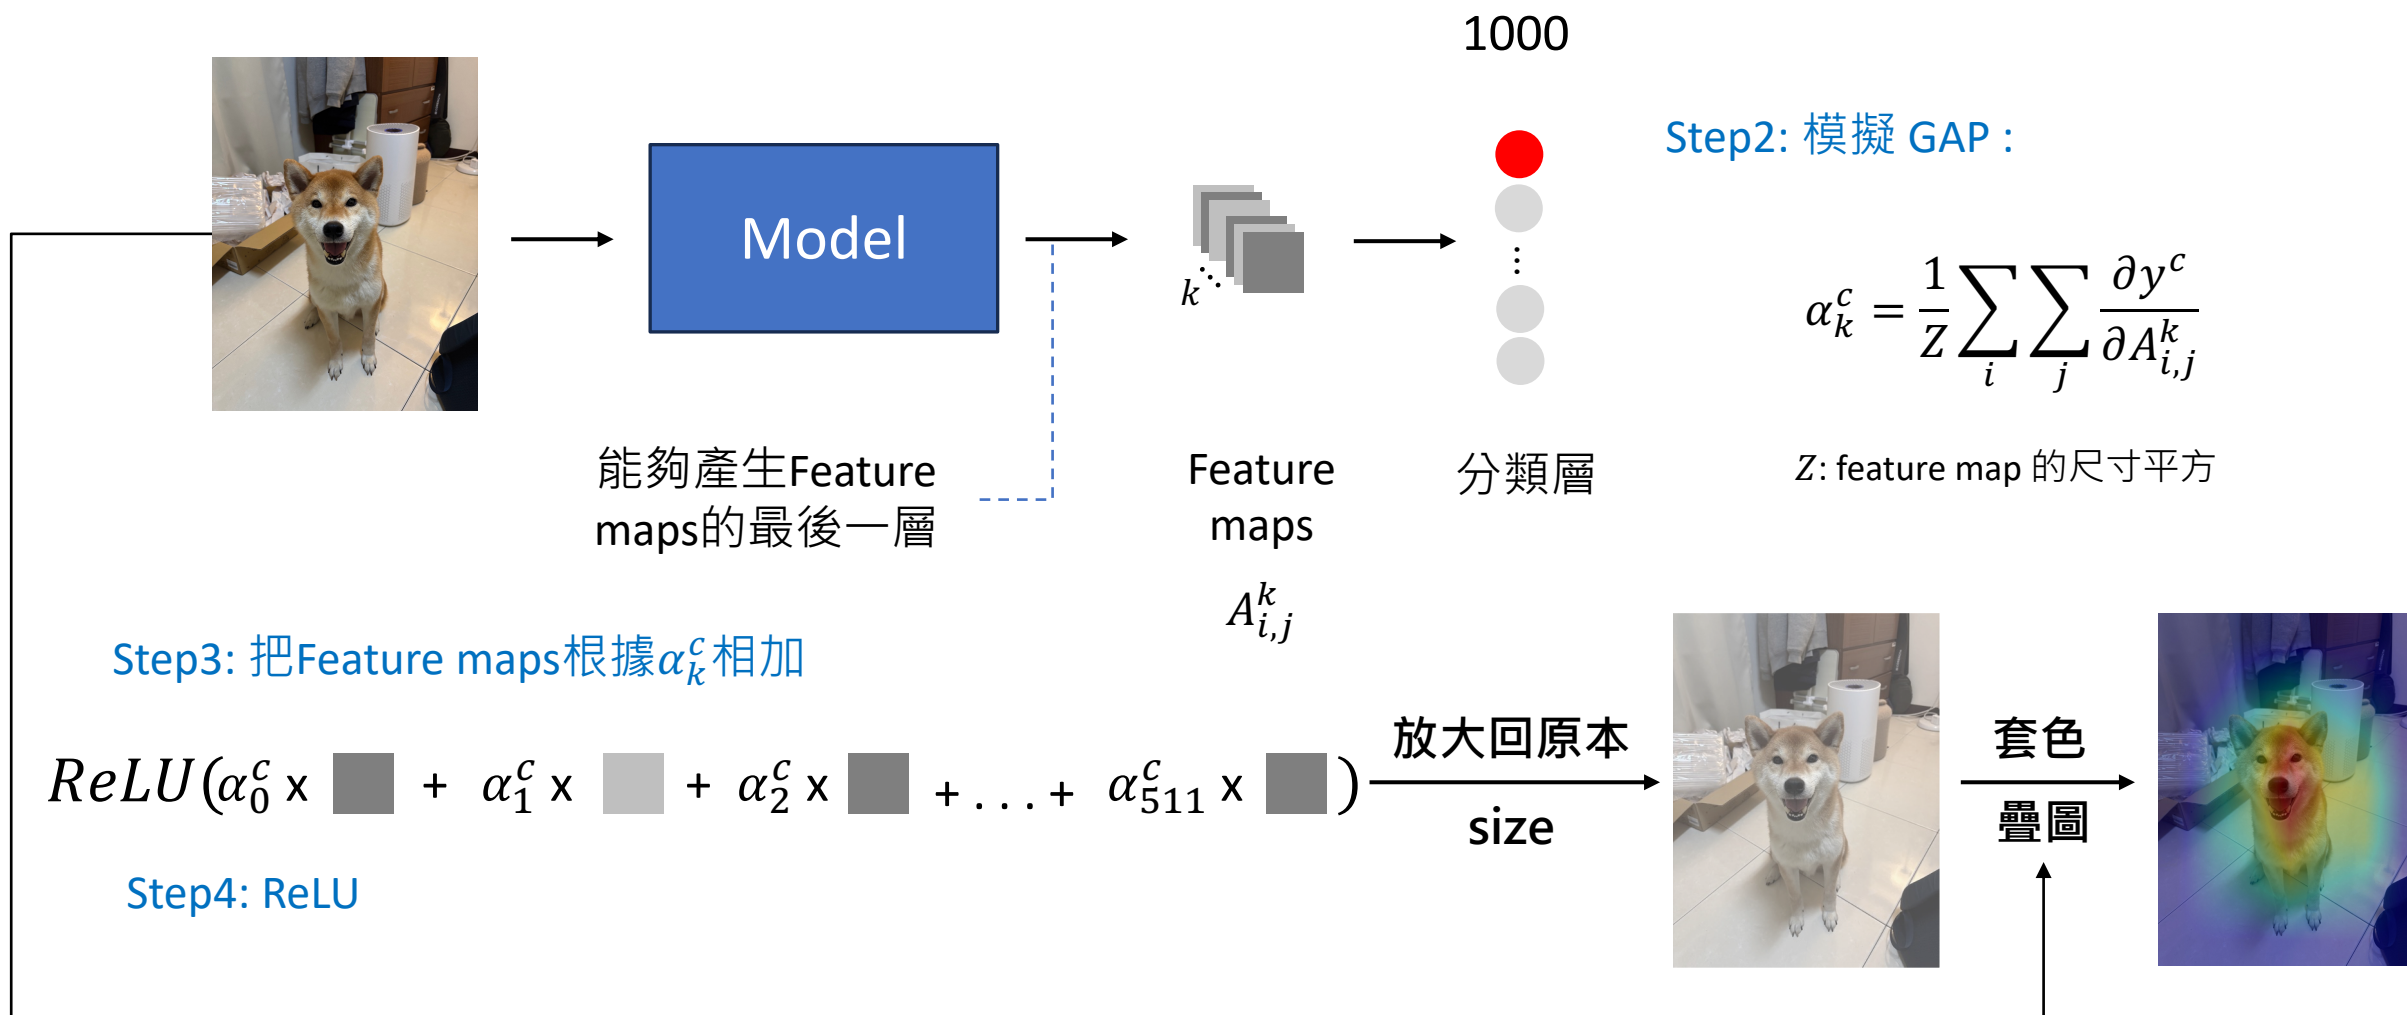
Step1: 計算特定 class (假設是 y^0)

對於任一張 feature map 中每個

位置 $(A_{i,j}^k)$ 的梯度 $\frac{\partial y^c}{\partial A_{i,j}^k}$

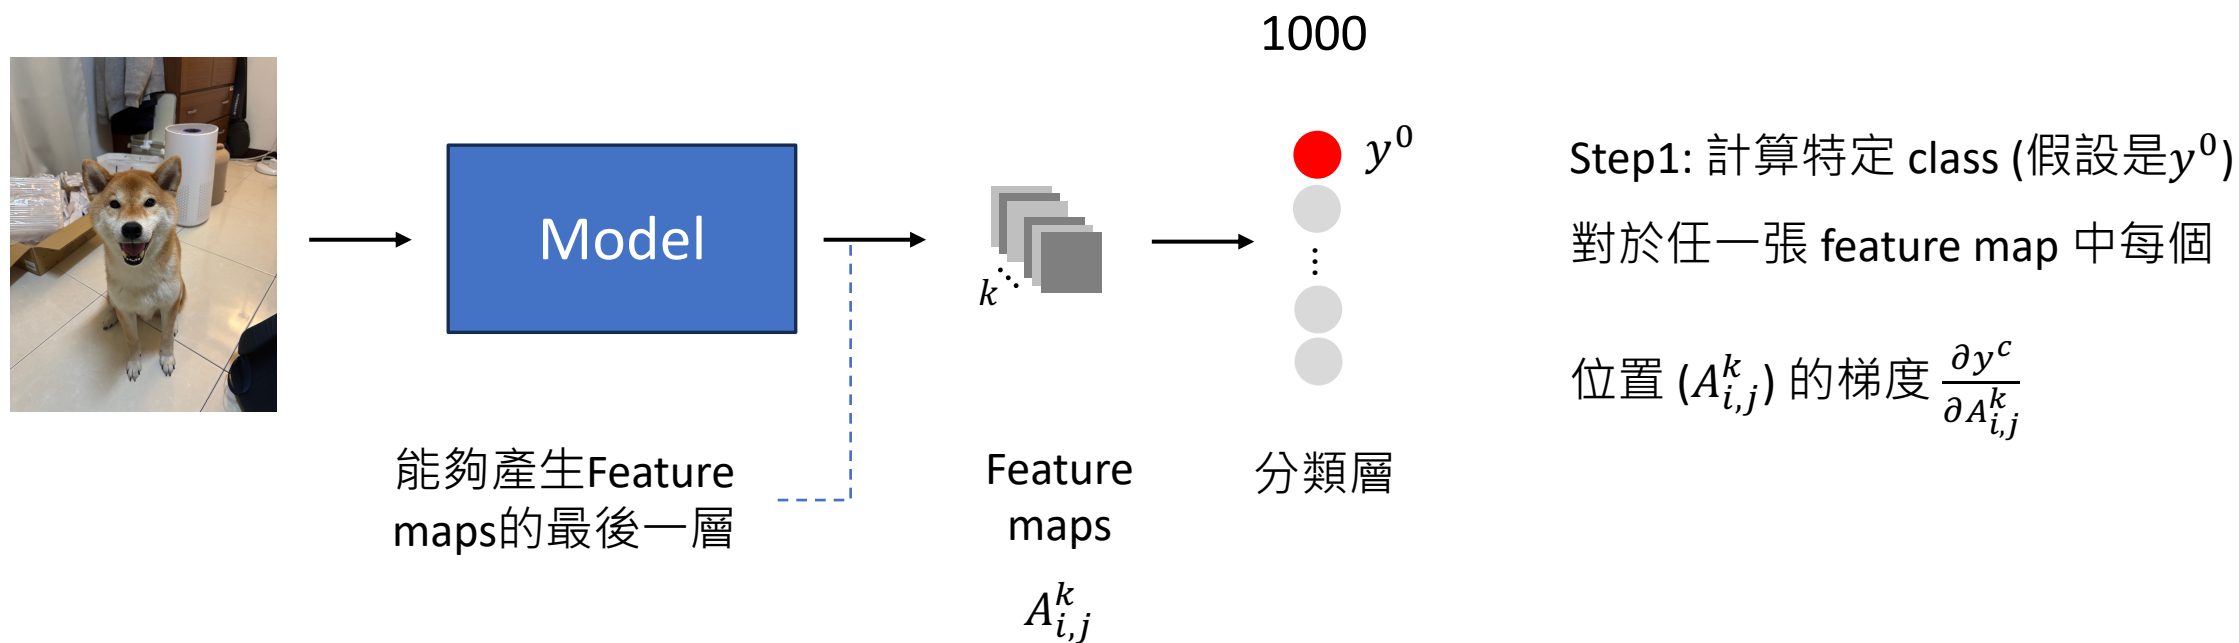


Grad-CAM (2)



為什麼要使用梯度？

i, j : feature map 中x軸與y軸位置
 A : feature map
 k : feature map 的數目
 y^c : 對應到 class c 的 label ID



- $\frac{\partial y^c}{\partial A_{i,j}^k}$ 代表 feature map 中任意位置的數值 ($A_{i,j}^k$) 對 y^c 的影響
- 如果一 feature map 有位置 $A_{i,j}^k$ 對 y^c 的影響很大 ($\frac{\partial y^c}{\partial A_{i,j}^k}$ 很大), α_k^c 也會跟著被放大, 代表該 feature map 可能對 y^c 特別重要



Automatic Evaluations

- Annotated datasets
 - ILSVRC (ImageNet Large Scale Visual Recognition Challenge) **Localization**
- Human Evaluation





Figure source: <https://www.kaggle.com/c/imagenet-object-localization-challenge/overview/description>



Example: Medical Image Classification

- Pelvic X-ray fracture classification with Grad-CAM



Example: VQA (Visual Question Answering)

- VQA with explanations

Whitehouse, Chenxi, Tillman Weyde, and Pranava Swaroop Madhyastha. "Towards a Unified Model for Generating Answers and Explanations in Visual Question Answering." Findings of EACL 2023.

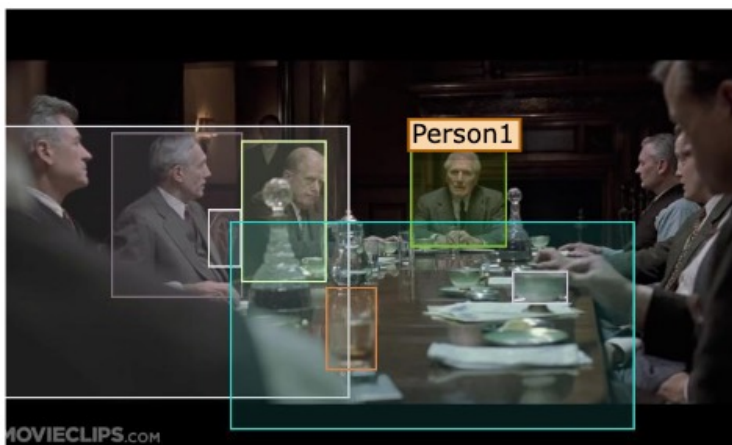


Question: What time of year was the picture likely taken?

Answer: fall

Ground Truth Explanations:

- 1) The child is wearing a long sleeve shirt and pants but no coat.
- 2) There are brown leaves on the sidewalk.
- 3) The time is fall.



Question: What is Person1 going to do?

Answer: Person1 is going to lead a business meeting.

Ground Truth Explanation:

Person1 is at the head of a table of men in suits.



重要論文

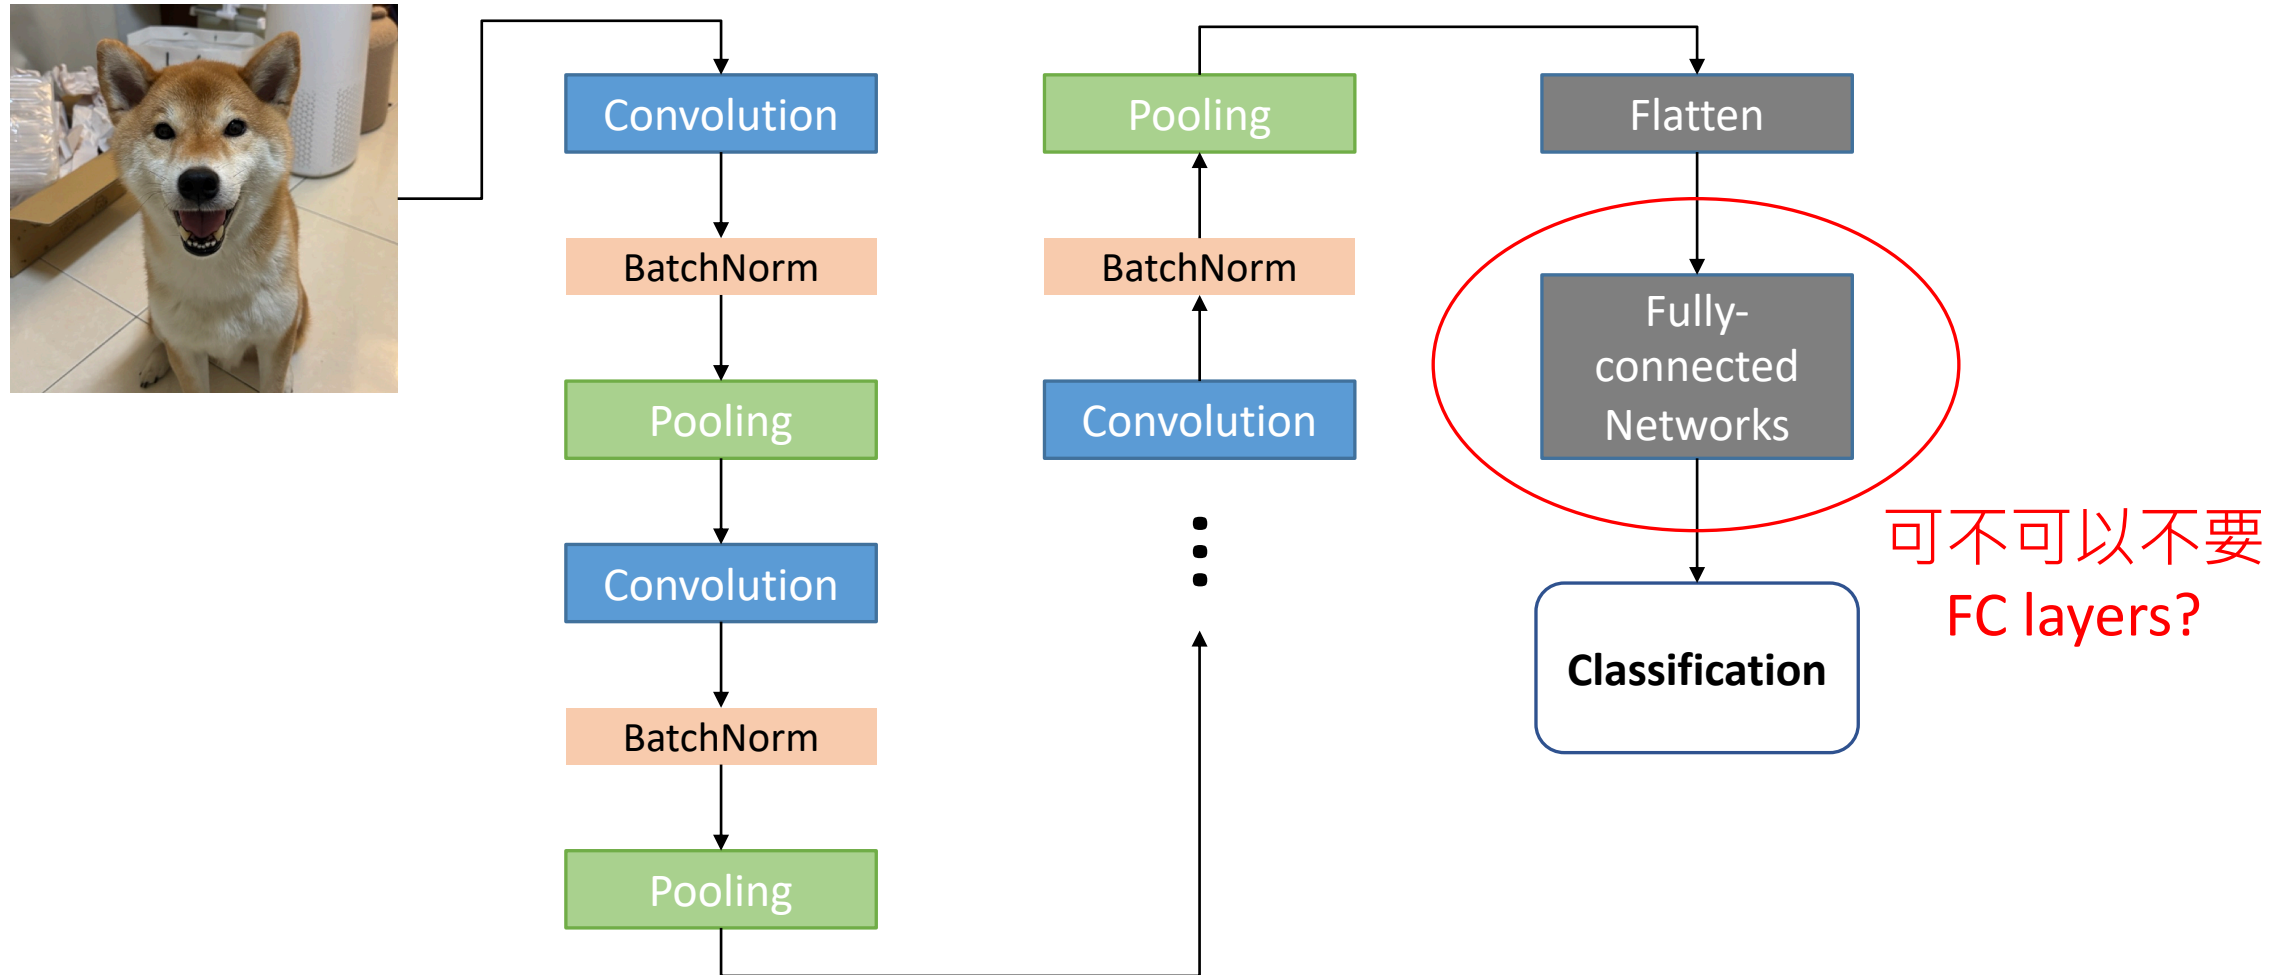
- Network In Network: <https://arxiv.org/abs/1312.4400>
- Class-activation Map (CAM): <https://arxiv.org/abs/1512.04150>
- Grad-CAM: <https://arxiv.org/abs/1610.02391>
- SHAP: <https://arxiv.org/abs/1705.07874>
- Guided back-propagation: <https://arxiv.org/pdf/1412.6806>



延伸主題

可不可以完全不要FC LAYERS?

FC layers 會大量增加參數



完全沒有採用 FC layers 的模型 (論文)


- Network In Network: <https://arxiv.org/abs/1312.4400>
- SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and <0.5MB model size (ICLR 2017): <https://openreview.net/forum?id=S1xh5sYgx>



Channels 數目調整

*維度意義：(C, H, W)

(128, 4, 4) (10, 4, 4)




PyTorch 寫法

```
torch.nn.Conv2d(  
    in_channels=128,  
    out_channels=10,  
    kernel_size=3,  
    padding=1,  
)
```

Filters 參數 (weights) 數量

$$128 * 10 * 3 * 3 = 11,520$$

(128, 4, 4) (10, 4, 4)



1 x 1 convolution

```
torch.nn.Conv2d(  
    in_channels=128,  
    out_channels=10,  
    kernel_size=1,  
    padding=0,  
)
```

$$128 * 10 * 1 * 1 = 1,280$$



1 x 1 Convolution

Stride = 1

1	1	1	1	0	0
0	1	1	0	1	0
0	0	1	1	0	0
0	0	1	1	1	0
0	0	0	1	1	0
0	0	0	0	0	0

1

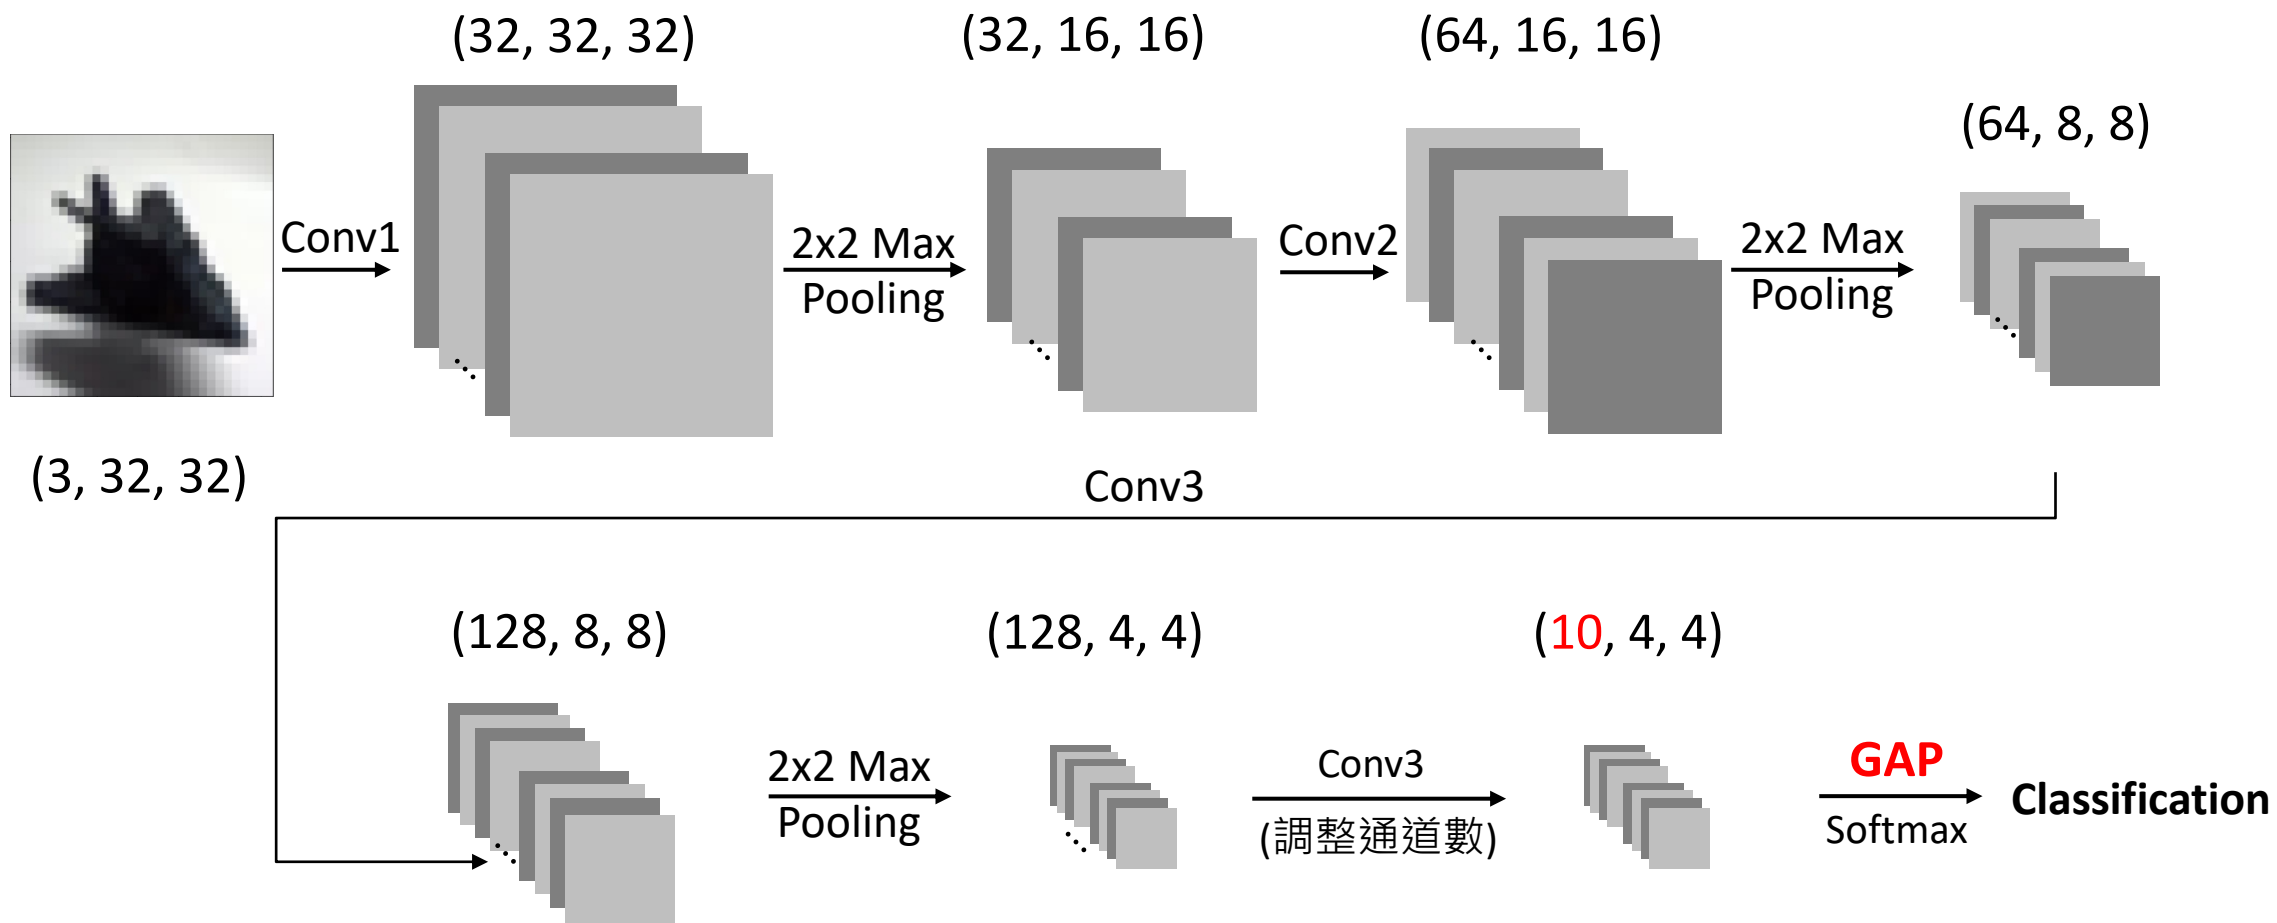
Filter
(假設數值為1)

1 1 1 1 0 0



CNN with Global Average Pooling (GAP)

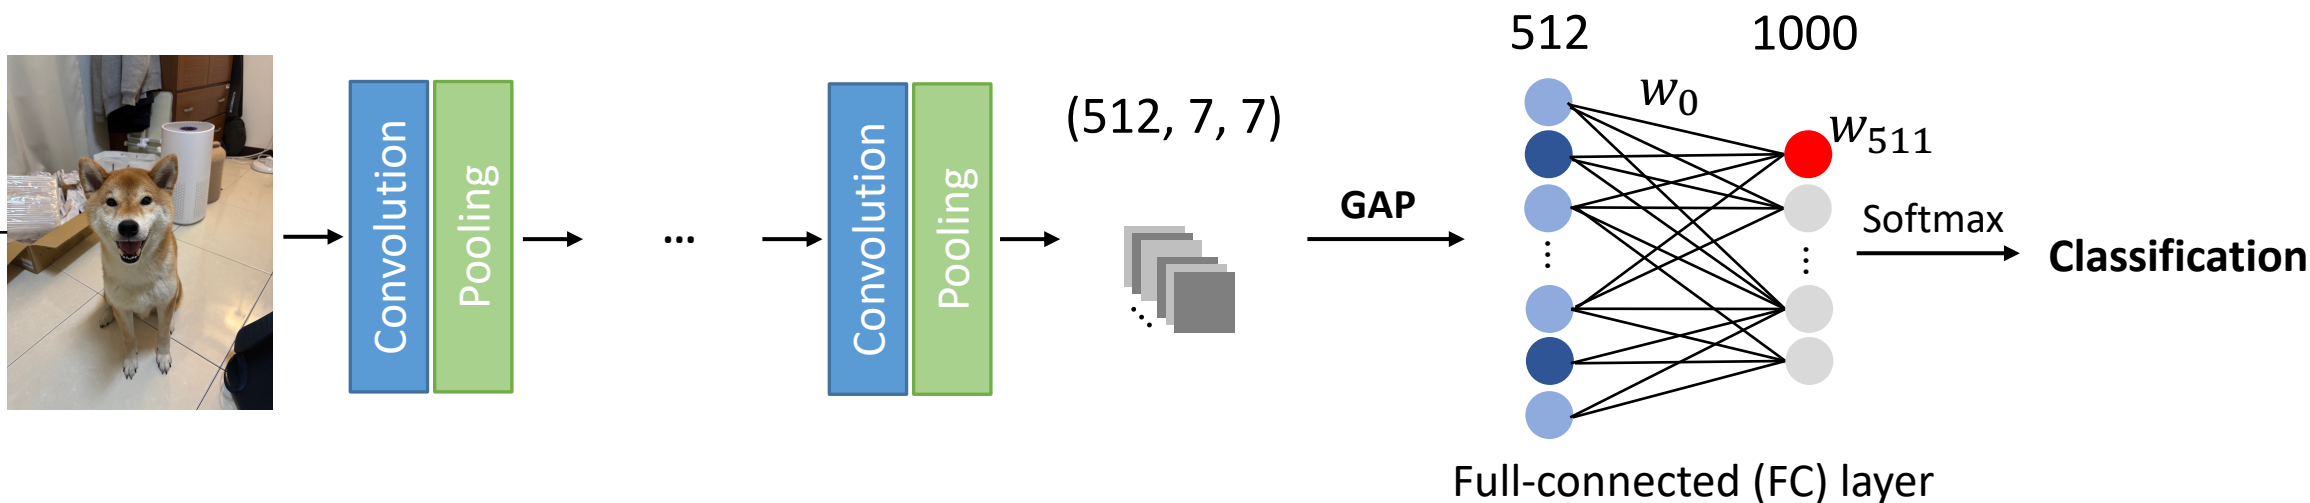
*維度意義：(C, H, W)



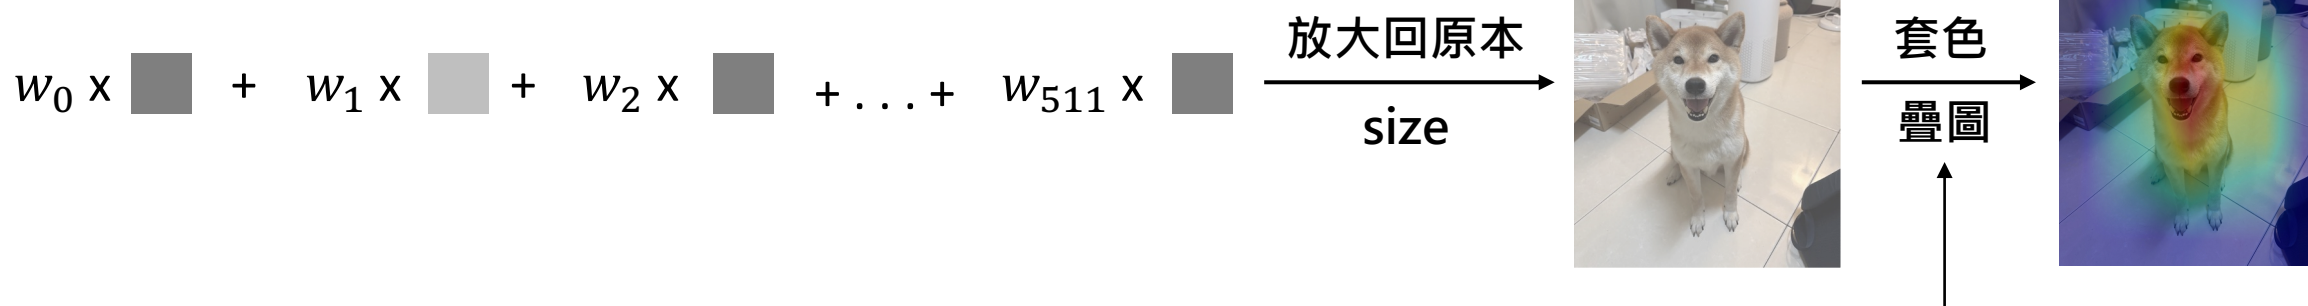
PyTorch

Class-activation Map (CAM)

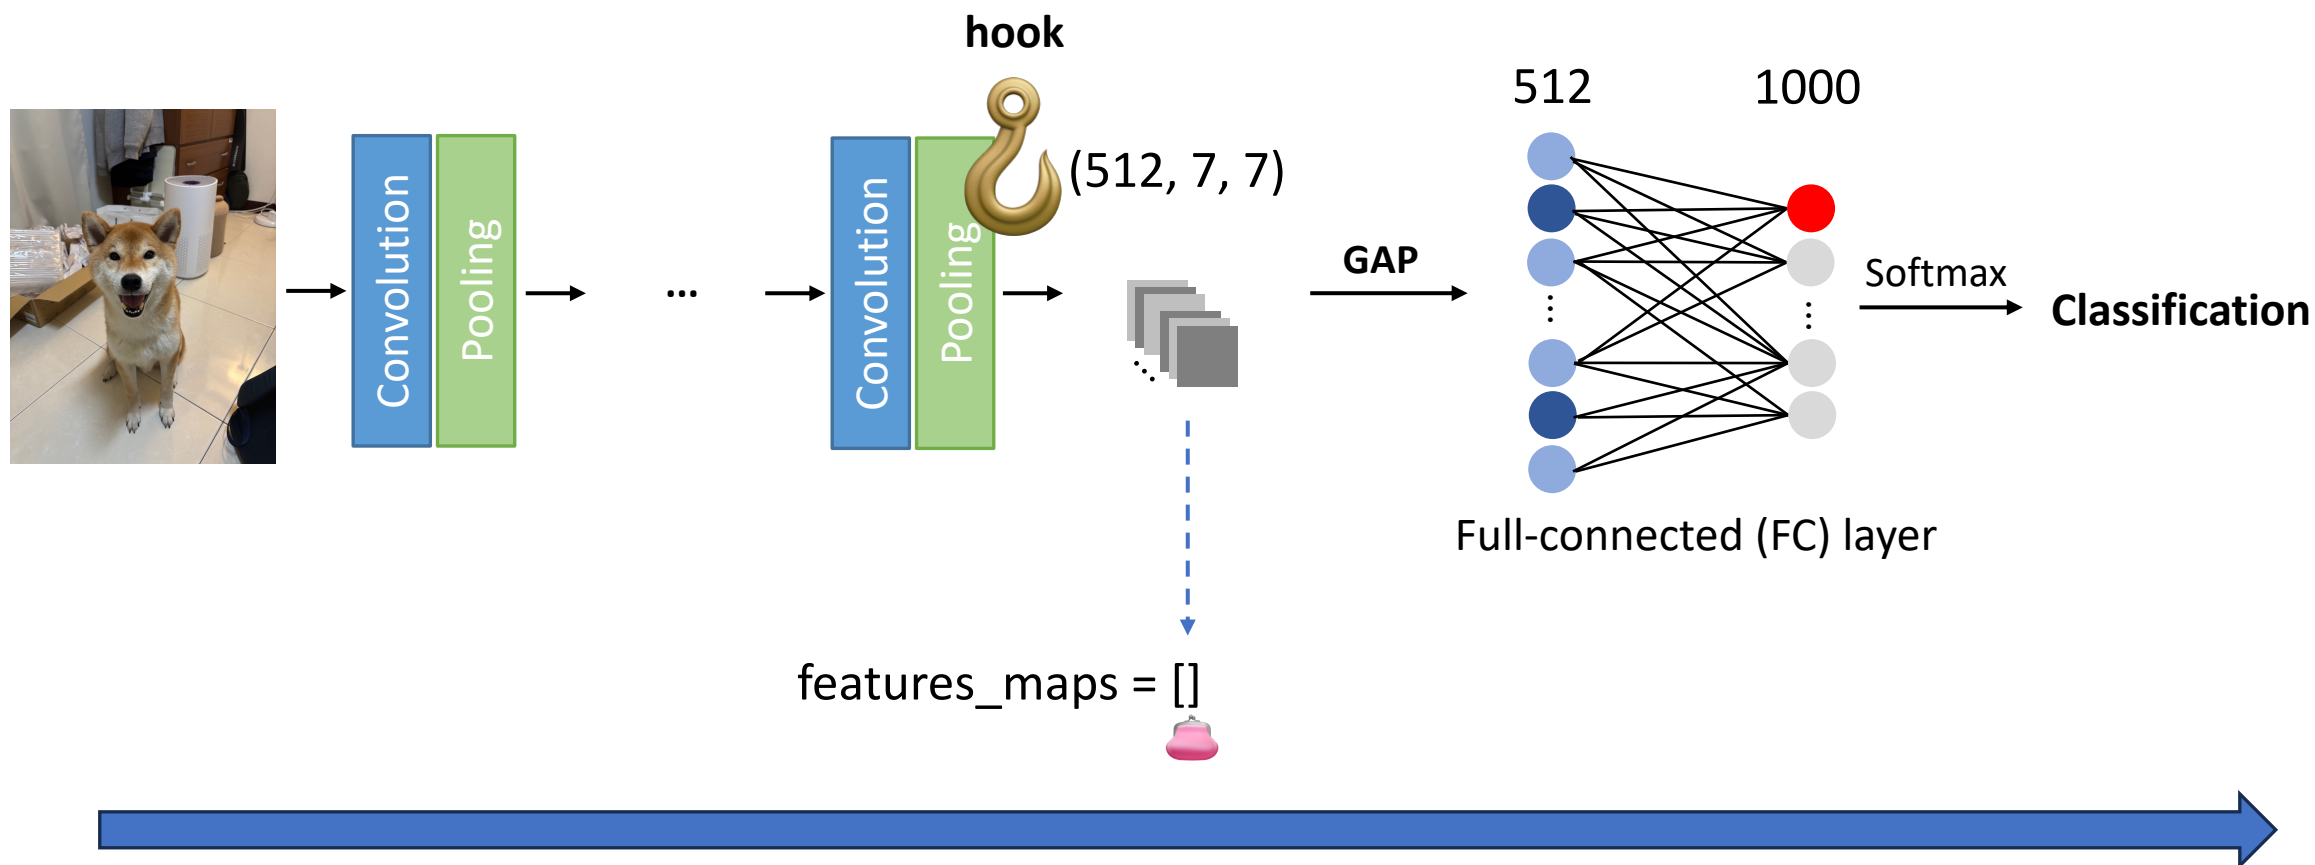
*假設使用 ImageNet pre-trained model



需要取得中間輸出



register_forward_hook



Thank you!

Instructor: 林英嘉

 yjlin@cgu.edu.tw

TA: 林君襄

 becky890926@gmail.com